

Word Embeddings with Word2vec

ADVISOR: PROF. AMITABHA MUKERJEE

PRANJAL SINGH

10327511

Index

Part I

1. word2vec by Mikolov et al.(2013)
2. Dataset
3. Results (English and Hindi)

Part II (Sentiment Analysis)

1. Tf-idf
2. Dataset
3. Results (English and Hindi)

PART I

word2vec


- Two architectures
 1. Continuous Bag-of-Word (CBOW)
 2. Skip Gram

CBOW: Predict the word given the context

Skip-Gram: Predict the context given the word

One Word Context (CBOW)

- Vocabulary size is V and hidden layer size is N
- Input vector is one-hot encoded vector, i.e., only one node of $\{x_1, \dots, x_V\}$ is 1 and others 0
- Weights between the input layer and the output layer can be represented by a $V \times N$ matrix \mathbf{W}




One Word Context (CBOW)

$$\mathbf{h} = \mathbf{x}^T \mathbf{W} = \mathbf{v}_{w_i}$$

\mathbf{v}_{w_i} is the vector representation of the input word w_i

$$u_j = \mathbf{v}'_{w_j}{}^T \cdot \mathbf{h}$$

u_j is score of each word in vocabulary and \mathbf{v}'_{w_j} is the j -th column of matrix \mathbf{W}'



One Word Context (CBOW)

We then use soft-max, a log-linear classification model, to obtain the posterior distribution of words, which is a multinomial distribution

$$p(w_j|w_I) = y_j = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})}$$

$$p(w_j|w_I) = \frac{\exp(\mathbf{v}'_{w_O}{}^T \mathbf{v}_{w_I})}{\sum_{j'=1}^V \exp(\mathbf{v}'_{w'_j}{}^T \mathbf{v}_{w_I})}$$

One Word Context (CBOW)

Update Equation for hidden->output weights

The training objective is to maximize the conditional probability of observing the actual output word w_o (denote its index in the output layer as j^*) given the input context word w_i with regard to the weights

$$\begin{aligned}\max p(w_o|w_i) &= \max y_{j^*} \\ &= \max \log y_{j^*} \\ &= u_{j^*} - \log \sum_{j'=1}^V \exp(u_{j'}) := -E\end{aligned}$$

Using Stochastic Gradient Descent, update equation becomes,

$$\mathbf{v}'_{w_j}{}^{(\text{new})} = \mathbf{v}'_{w_j}{}^{(\text{old})} - \eta \cdot e_j \cdot \mathbf{h} \quad \text{for } j = 1, 2, \dots, V.$$

One Word Context (CBOW)


Update Equation for input->hidden weights

$$\mathbf{v}_{w_I}^{(\text{new})} = \mathbf{v}_{w_I}^{(\text{old})} - \eta \cdot \text{EH}$$

Multi Word Context (CBOW)

$$\begin{aligned}\mathbf{h} &= \frac{1}{C} \mathbf{W} \cdot (\mathbf{x}_1 + \mathbf{x}_2 + \cdots + \mathbf{x}_C) \\ &= \frac{1}{C} \cdot (\mathbf{v}_{w_1} + \mathbf{v}_{w_2} + \cdots + \mathbf{v}_{w_C})\end{aligned}$$

$$\begin{aligned}E &= -\log p(w_O | w_{I,1}, \cdots, w_{I,C}) \\ &= -u_{j^*} + \log \sum_{j'=1}^V \exp(u_{j'}) \\ &= -\mathbf{v}'_{w_O} \cdot \mathbf{h} + \log \sum_{j'=1}^V \exp(\mathbf{v}'_{w_j} \cdot \mathbf{h})\end{aligned}$$



Multi Word Context (CBOW)

Update Equations


$$1) \quad \mathbf{v}'_{w_j}{}^{(\text{new})} = \mathbf{v}'_{w_j}{}^{(\text{old})} - \eta \cdot e_j \cdot \mathbf{h} \quad \text{for } j = 1, 2, \dots, V.$$

$$2) \quad \mathbf{v}_{w_{I,c}}{}^{(\text{new})} = \mathbf{v}_{w_{I,c}}{}^{(\text{old})} - \frac{1}{C} \cdot \eta \cdot \mathbf{EH} \quad \text{for } c = 1, 2, \dots, C.$$

Skip-Gram Model

- It is the opposite of CBOW Model
- On the output layer, instead of outputting one multinomial distribution, we are outputting C multinomial distributions.
- Each output is computed using the same hidden \rightarrow output matrix


$$p(w_{c,j} = w_{O,c} | w_I) = y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})}$$



Skip-Gram Model

Loss Function

$$\begin{aligned} E &= -\log p(w_{O,1}, w_{O,2}, \dots, w_{O,C} | w_I) \\ &= -\log \prod_{c=1}^C \frac{\exp(u_{c,j_c^*})}{\sum_{j'=1}^V \exp(u_{j'})} \\ &= -\sum_{c=1}^C u_{j_c^*} + C \cdot \log \sum_{j'=1}^V \exp(u_{j'}) \end{aligned}$$



Optimization

1. Hierarchical Softmax (Morin & Bengio, 2005)
2. Negative Sampling (Gutmann & Hyvarinen, 2012)

Dataset

- Hindi :Wikipedia text dump (290 MB) ~4 mins of training
 - 23848940 words
 - 723737 words in vocabulary
 - 106876 words after removing words with *freq*<5
- English: Wikipedia text dump (9.3 GB) ~3.5 hrs of training
 - 1,703,849,452 words
 - 13,027,758 words in vocabulary
 - 1,778,685 words after removing words with *freq*<5

Result (English)

Odd One out

1) breakfast cereal dinner lunch

Ans: Cereal

2) eight six seven five three owe nine

Ans: owe

3) math shopping reading science

Ans: shopping

Result (English)

Top Similar Words

Father

- 1) grandfather
- 2) uncle
- 3) mother
- 4) father-in-law
- 5) brother

FRANCE

- 1) GERMANY
- 2) FRENCH
- 3) GREECE
- 4) NETHERLANDS
- 5) SCOTLAND

XBOX

- 1) XBLA
- 2) Xbox360
- 3) SmartGlass
- 4) 360/PS3
- 5) Playstation
- 6) Qubed
- 7) Kinect

scratched

- 1) scraped
- 2) rubbed
- 3) bruised
- 4) cracked
- 5) discarded
- 6) shoved
- 7) tripped
- 8) scratches
- 9) ripped
- 10) pasted

megabits

- 1) gigabits
- 2) kilobits
- 3) megabit
- 4) terabits
- 5) MB/s
- 6) Tbit/s

Result (Hindi)

Odd One out

1) भारत रूस मुम्बई चीन

Ans: मुम्बई

2) लड़की बेहन महिला मर्द

Ans: मर्द

3) नेता मंत्री सरकार उद्योग

Ans: उद्योग

Result (Hindi)

Top Similar Words

भारत

- 1) प्रदेश
- 2) तिब्बत
- 3) देश
- 4) आँध्रप्रदेश
- 5) लद्दाख

व्यापार

- 1) व्यवसाय
- 2) पुनर्बीमा
- 3) वाणिज्य
- 4) बैंकिंग
- 5) उद्योग

ओबामा

- 1) क्लिंटन
- 2) बराक
- 3) सीनेटर
- 4) राष्ट्रपति
- 5) उम्मीदवार

PART II

Sentiment Analysis

- Sentiment analysis aims to determine the attitude of a speaker or a writer with respect to some topic or the overall contextual polarity of a document.
- A basic task in sentiment analysis is classifying the *polarity* of a given text at the document, sentence, or feature/aspect level — whether the expressed opinion in a document, a sentence or an entity feature/aspect is positive, negative, or neutral.

(Source: Wikipedia)

tf-idf (term frequency–inverse document frequency)

- It is a numerical statistic that is intended to reflect how important a word is to a document in a collection or corpus
- The tf-idf value increases proportionally to the number of times a word appears in the document, but is offset by the frequency of the word in the corpus

$$tfidf(d, t) = tf(d, t) \times \log\left(\frac{|D|}{df(t)}\right)$$

Here $df(t)$ is the number of documents in which term t appears.

D is the total number of documents

$tf(d, t)$ is the term-frequency

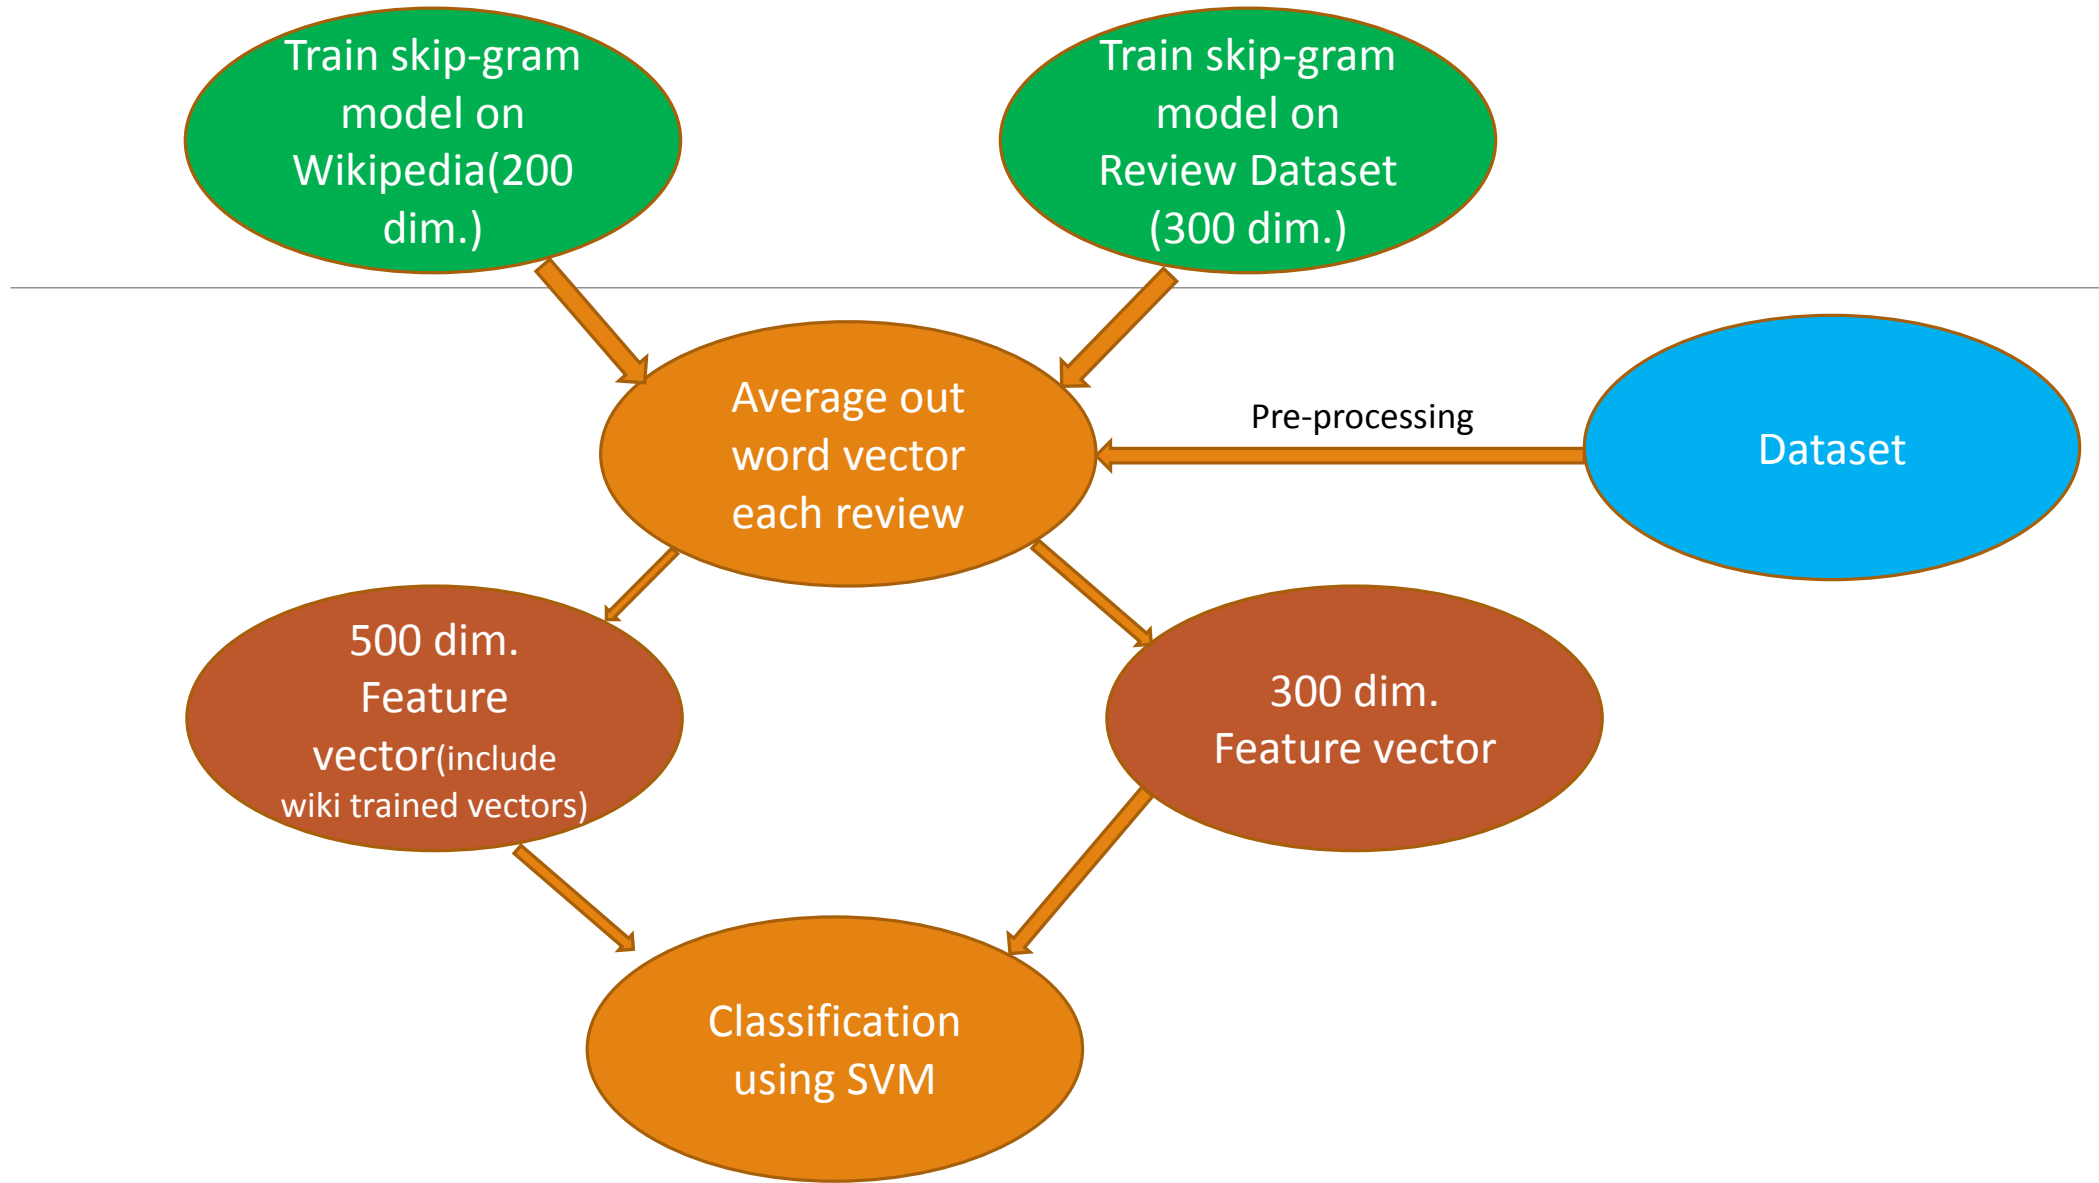
Dataset

1. IMDB 50,000 Movie Review Dataset(English)
 - I. Contains 25,000 Positive and 25,000 Negative reviews
 - II. 25,000 training examples and 25,000 testing examples

2. Hindi Product Review Dataset(IIT)
3. Hindi Movie Review Dataset(IITB)

Methodology

- Trained word2vec on each dataset for respective evaluation
- Features are constructed by taking average of word vectors for each review(not weighted average)
- Ignored stop-words(the, a, there, etc.) in English but not Hindi(मैं, वह, ने, etc.)
- Also constructed feature set by merging word vectors of Wikipedia trained word2vec(Hindi & English Wikipedia)
 - E.g. 300 dim.(from first training) + 200 dim.(from second training)= 500 dim. Feature vector
- Used SVM implementation of libSVM in scikit



Result(English-IMDB)

Classifier	Accuracy(%)
SVM(linear)	87.05
Random Forest(280 Trees)	84.14
Naïve Bayes	75.95
Logistic Regression(l1 penalty)	86.90
k-NN(20)	76.76
SVM(linear) with Wiki trained word vectors merged(300+200 dim)	87.56

*** Skip-gram model was trained with vector dimension as 300 and min word-count as 20**

Result(Hindi-Product Review: IIIT)

Feature Set	Accuracy(%)
500 dimensional word vector	78.0
500 dimensional word vector + tf-idf	90.73
Using subjective Lexicon construction(Thesis by Piyush Arora, IIIT'13)	79.62

- **Skip-gram model was trained with vector dimension as 300 and min word-count as 20**
- **Used Linear SVM as classifier**

Result(Hindi-Movie Review: IITB)

Feature Set	Accuracy(%)
500 dimensional word vector	79.62
500 dimensional word vector + tf-idf	88.13
Tf-idf & n-gram ^[6]	78.14

- Skip-gram model was trained with vector dimension as 300 and min word-count as 20
- Used Linear SVM as classifier

References

1. Manning, W. Y. (2013). Bilingual Word Embeddings for Phrase-Based Machine Translation. *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*, (p. 2013).
2. Weston, R. C. (2011). Natural Language Processing (Almost) from Scratch. *Journal of Machine Learning Research*, 12, 2493--2537.
3. Mikolov, T., Chen, K., Corrado, G. & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. *CoRR*, abs/1301.3781.
4. Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S. & Dean, J. (2013). Distributed Representations of Words and Phrases and their Compositionality. In C. Burges, L. Bottou, M. Welling, Z. Ghahramani & K. Weinberger (ed.), *Advances in Neural Information Processing Systems 26* (pp. 3111--3119) .
5. Aditya Joshi, Balamurali A. R. and Pushpak Bhattacharyya. A Fall-back Strategy for Sentiment Analysis in Hindi: a Case Study. International Conference on Natural language Processing (ICON), Karagpur.-2010
6. Sarath Chandar A P, Mitesh M Khapra, Ravindran B, Vikas Raykar, Amrita Saha, "Multilingual Deep Learning". In Deep Learning Workshop at NIPS 2013.

Thank You!!!

NLP from Scratch

- Built a unified architecture for tasks such as POS tagging, Chunking, NER
- Compared against classical NLP benchmarks
- Avoided task specific engineering
- Generalize a system to handle multiple tasks

NLP from Scratch

- Learn lookup table by back propagation
- Words are mapped to *d-dimensional* vector using lookup table operation
- Lookup table returns a matrix for a given sentence

NLP from Scratch

Used entire English Wikipedia to learn word embeddings (631 million words)

Tokenized using Penn Treebank Tokenizer

The total training time was about four weeks

Window size: 11 and a Hidden layer with 100 units

Objective: Seek a network that computes a higher score when given a legal phrase than when given an incorrect phrase

$$\theta \mapsto \sum_{x \in \mathcal{X}} \sum_{w \in \mathcal{D}} \max \left\{ 0, 1 - f_{\theta}(x) + f_{\theta}(x^{(w)}) \right\}$$