

Decompositional Semantics for Improved Language Models

*A Thesis Submitted
in Partial Fulfilment of the Requirements
for the Degree of*

Master of Technology

by

Pranjal Singh

Roll No.: 10327511

under the guidance of
Dr. Amitabha Mukerjee



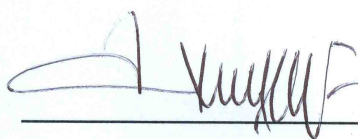
Department of Computer Science and Engineering

Indian Institute of Technology Kanpur

June, 2015

CERTIFICATE

It is certified that the work contained in this thesis entitled “*Domain Driven Decompositional Semantics*”, by *Pranjal Singh* (Roll No. 10327511), has been carried out under my supervision and that this work has not been submitted elsewhere for a degree.



(Dr. Amitabha Mukerjee)

Department of Computer Science and Engineering,
Indian Institute of Technology Kanpur
Kanpur - 208016

May, 2015

Abstract

In recent years, distributional semantics or vector models for words have been proposed to capture both the syntactic and semantic similarities between words. Such vectors may be obtained for words as used in a large corpus, or in a given domain. Since these are language free models and can be obtained in an unsupervised manner, they are of interest for under-resourced languages such as Hindi. We start with an overview which shows that a reasonable measure of semantic similarity in Hindi seems to be captured by a word vector map.

A difficulty with word vectors is that of mapping larger compositions. Several methods - additive, multiplicative, or tensor neural models, have been proposed. Other methods have exploited logarithmic time complexity to learn paragraph vectors on the fly. In most such methods all words are treated equally (except for a few classes like stop words). Here we propose a simple innovation based on graded weights for the words in compositional texts. We experiment with different weighting schemes and find that using the well-known tf-idf weights outperforms most of the standard approaches. On the extensively-researched benchmark task of sentiment analysis, our method achieves an accuracy of 94.19% on the IMDB movie review dataset (improvement of 1.6% over previous best). Even greater improvements are shown for the Amazon product review datasets. We also implement an ensemble model which boosts our accuracy by about 0.5% using recurrent neural network. Finally, we also consider the language free aspects by demonstrating an increment of over 12% on existing results for Hindi.

Dedicated to my Parents, Brother and Friends for their love and support

Acknowledgement

I would like to take this opportunity to thank the people who have helped me in preparation of this thesis.

First of all, I would like to express my sincere gratitude towards my thesis supervisor, Dr. Amitabha Mukerjee, for his constant support and encouragement. I am grateful for his patient guidance and advice which gave a proper direction to my efforts. I am grateful to him for providing me ample freedom to choose a topic of my interest and deciding how to pursue it.

I thank the Department of Computer Science and Engineering, IIT Kanpur, for providing the necessary infrastructure and a congenial environment for research work. I thank Tomas Mikolov for making the code of word2vec and sentence2vec publicly available.

I am also thankful to all my friends and specially wingmates (chaman, ladha, chotu, anshu, lassu, mm, baba, aryan, deo and few more) who have made my stay at IIT Kanpur most enjoyable and who in general gave me strength and hope whenever I required. Last, but not the least, I thank my parents for always being there for me.

Pranjal Singh

Contents

Abstract	iv
List of Tables	xiv
List of Figures	xvi
List of Algorithms	xvii
1 Introduction	1
1.1 The Problem	6
1.2 Contributions of this Thesis	6
1.3 Organization of this Thesis	7
2 Related Work and Background	9
2.1 Word-Vectors and Representations	11
2.2 Document Vectors	15
2.3 Recurrent Neural Network	17
2.4 Semantic Composition	20
3 Data Acquisition	23
3.1 Hindi Corpus	23
3.2 English Corpus	25
4 Dimensionality Reduction	27
4.1 Principal Component Analysis(PCA)	28
4.2 ANOVA: F-value	29

4.3	Result	30
5	Models and Experiments	33
5.1	Context Size	33
5.2	SkipGram or CBOW	35
5.3	Skip-Gram and <i>tf-idf</i> based Word Vectors	36
5.4	Weighted Average Word Vectors	37
5.5	Word and Document Vectors with tf-idf: Enhanced Document Vectors	39
5.6	Ensemble of RNNLM and Enhanced Document Vectors	39
5.7	Embeddings	40
5.8	Work Flow	42
5.9	Technical Specifications	43
6	Results	45
6.1	Results on English Datasets	45
6.2	Results on Hindi Datasets	48
6.3	Odd One Out	49
6.4	Similar Words	50
7	Conclusion and Future Work	51
7.1	Conclusion	51
7.2	Future Work	53
7.3	Summary	54
	Bibliography	55

List of Tables

2.1	Comparison of WSJ results obtained with various models(RNN is trained on just 6.4M words)	19
2.2	Results of Vector Composition with different Operations	20
3.1	Statistics of Dainik Jagran Movie Reviews	24
3.2	Statistics of Navbharat Times Movie Reviews	24
3.3	Statistics of Movie Reviews from Jagran and Navbharat	24
4.1	Accuracies on 700-Movie Review Dataset	30
4.2	Accuracies on IITB Movie Review and IIIT Product Review Dataset using Document Vectors	31
5.1	Results on IMDB Movie Reviews using Various Weighting Techniques(Composite Document Vector);Accuracy(2) is when we exclude tf-idf features . . .	39
6.1	Comparison of accuracies on 3 Datasets(IMDB, Amazon Electronics Review and Hindi Movie Reviews(IITB)) for various types of document composition models. The state of the art for these tasks are: IMDB: 92.58% [12]; Amazon:85.90% [6], Hindi:79.0% [1].	45
6.2	Results on IMDB Movie Review Dataset	45
6.3	Results on IMDB Movie Review Dataset	46
6.4	Comparison of results on IMDB Movie Review Dataset with Various Features	46
6.5	Results on Amazon Electronics Review Dataset	47
6.6	Accuracies for Product Review and Movie Review Datasets.	48

6.7	Comparison of Approaches: Product Review Dataset	48
6.8	Comparison of Approaches: Movie Review Dataset	48
6.9	Some sentiment words and their neighbors	49
6.10	Odd One Out in English	49
6.11	Odd One Out in Hindi	49
6.12	Top Few Similar words in English	50
6.13	Top Few Similar words in Hindi	50

List of Figures

2.1	Skip Gram Model(Figure from Rong (2014))	13
2.2	Framework for learning word vectors(Figure from Le (2014)).	16
2.3	Framework for learning paragraph vectors(Figure from Le (2014)).	17
2.4	Simple Recurrent Neural Network(Figure from [18]).	18
2.5	Approach of Recursive Neural Network(Figure from [31]).	21
4.1	PCA of a multivariate Gaussian distribution centered at (1,3) with a standard deviation of 3 in roughly the (0.878, 0.478) direction and of 1 in the orthogonal direction. The vectors shown are the eigenvectors of the covariance matrix scaled by the square root of the corresponding eigenvalue, and shifted so their tails are at the mean(Figure from [36]).	29
5.1	High Dimensional representation of Wiki Text with a) Context Size 5 b) Context Size 10	34
5.2	Variation of Accuracy with Different Context Size on Watches and MP3 Datasets.	35
5.3	Variation of Accuracy with skipgram and cbow on Watches and MP3 Datasets.	36
5.4	t-SNE visualization of the top 5000 Hindi Words in high dimensional space. (Magnify to see details).	40
5.5	A closer look at two clusters in the visualization showing a) quantity relations, b) locations and c) diseases.	41
5.6	Work Flow Summary.	42

6.1	Accuracies of Different Classifiers with Average Word Vectors(IMDB).	
	47

List of Algorithms

1	PCA Algorithm	28
2	Weighted Average for Vector Composition	37

Chapter 1

Introduction

Language representation is a very crucial aspect to perform various NLP tasks and has been looked into great detail in recent times. Language representation models have fallen into broadly two categories: ones which require hand-trained language databases such as treebanks (e.g.,[31]) and ones which are language agnostic and work on raw corpora (e.g., LDA, BOW, SkipGram, NLM, etc.). Liu [13] compare various language agnostic models for topic modeling. We have tried to present a work which is language independent and which is enriched with vital properties of other models. This ensures that our model is effective as well as accurate in various languages. The world class results in English clearly indicate the efficacy of our approach and improvements in Hindi depict the deficiency in other models which were used earlier.

The large decrease in computational complexity of recent models [20, 17, 12] has led us to experiment with various factors and parameters helping us in providing justification and validation to arguments which were not experimentally justified but claimed on the basis of intuition.

Language independent models such as LDA and BOW have been quite effective since long time. Variants of BOW such as tf-idf had changed the perception of researchers towards these models when they were proved effective in various NLP tasks. LDA was able to model inter and intra documental statistical and relational structure quite well overcoming the drawbacks of BOW. But the semantic and syn-

tactical dependencies were still ignored. After the introduction of neural language vector models, NLP saw a huge diversion in representation of words and documents.

Vector models for individual words are obtained via distributional learning, the mechanisms for which varies from document-term matrix factorization [10], various forms of deep learning [5, 32, 31], optimizing models to explain co-occurrence constraints [20, 27], etc. Once the word vectors have been assigned, similarity between words can be captured via cosine distances. The same models have been extended [12] with new variables to build vector models for sentences and documents. These models include the essence of individual words as well as their relative order in terms of sentence vector which was earlier absent in word vectors. The advantage of these approaches is that they can capture both the syntactic and the semantic similarity between words/documents in terms of their projections onto a high-dimensional vector space; further, it seems that one can tune the privileging of syntax over semantics by using local as opposed to large contexts [8].

There has been another problem which has attracted the attention of Indian grammarians for over a period of nearly a millennium. They have been trying to find whether sentence meaning accrues by combining word meanings, or whether words gain their meanings based on the context they appear in [14]. The former position, that meaning is *compositional*, has been associated with the Fregean enterprise of semantics, whereas recent models, building on large corpora of text (and associated multimedia) a large degree of success has accrued to models that attempt to model word meaning based on their linguistic context (e.g. [10]). The latter line has resulted in strong improvements in several NLP tasks using word vectors [5, 32, 20, 31].

For resource-poor languages, these approaches have the added lure that many of these methods are completely unsupervised and work directly with large raw text corpora, thus avoiding contentious issues such as deciding on a POS-tagset, or expensive human annotated resources such as treebanks. For Indian languages which are highly inflected, stemming or identifying the lemma is another problem which

such models can overcome, provided the corpus is large enough. Nonetheless, this approach remains under-explored for Indian languages. At the same time, it must be noted that many approaches seek to improve their performance by combining POS-tags and even parse tree structures into the models for higher accuracies in specific tasks [31].

One problem in this approach is that of combining the word vectors into larger phrases. In past work, inverse-similarity weighted averaging appears to work to some extent even for complex tasks such as essay grading [11], but multiplicative models (based on a reducing the tensor products of the vectors) appears to correlate better with human judgements [22, 31]. Another complexity in composition is that composing words across phrasal boundaries are less meaningful than composing them within a phrase - this has led to models that evaluate the nodes of a parse tree, so that only coherent phrases are evaluated [31]. The results reported here, are based on applying the Skip-Gram model [17] to Hindi.

[33] give a detailed overview of various vector space models and their composition. A surprising event in Information Theory has higher information content than an expected event (Shanon, 1948). The same happens when we give weights to word vectors. We give more weights to events which occur by surprise and less weight to events which are expected. The most popular weighting concept in this domain is the idea of tf-idf which is explained later. [28] defined a large family of tf-idf weighting functions and evaluated them on information retrieval tasks, demonstrating that tf-idf weighting can yield significant improvements over raw frequency. Term weighting is also significant in reducing the weights of highly correlated words such as *hostage* and *hostages* when they occur together (Church, 1995). [4] have found use of term weighting in sentiment classification in Thai language.

Lastly, a single learning algorithm in any domain cannot always induce the most accurate learner and hence we include in this work, the concept of ensemble learning though it is not ensemble in absolute sense. We require that our classifier should have properties of both a discriminative model and a generative classifier and hence

merged the predictions of both achieving better results(majority voting). Since our main focus is not ensemble, so we donot look into other ensemble methods in detail.

1.0.1 Sentiment Analysis

In order to evaluate the efficacy of the model, we apply it to the task of sentiment analysis. Here the problem is that of identifying the polarity of sentences (Liu et al. 2012); for example:

- Positive: रामू ने कहानी की रफ़्तार कहीं थमने नहीं दी [Ramu didn't allow the pace of the story to subside]
- Negative: पर्दे पर दिखाया जा रहा खौफ़ सिनेमाघर मे नहीं पसर पाता [The horror shown on the screen didn't reach the theater]

This is a problem that has attracted reasonable attention in both Hindi and English (see section 2). There have been heuristic based and machine learning based models used in this domain [35]. Heuristic based methods, in general, classify text sentiments on the basis of total number of derived positive or negative sentiment oriented features. For example, Hatzivassiloglou and McKeown [7] described that adjectives are more predictive of sentiment classification and they predicted the sentiment of adjectives by inspecting them in conjunction with "and", "or", "but", "either/or" and "neither/nor". However the problem with this approach is that even if adjectives are important in many languages but this model underestimates some predictive words of other parts-of-speech which may be important in some languages. Also these models rely heavily on human engineered features which, in general, is a domain and language dependent task.

In this work, we first learn a distributional word vector model as well as sentence vector model based on the wikipedia corpus as well as the sentiment corpus, and then we use this to discern the polarities on the existing corpora of movie and product reviews. To our own surprise, we find that even a simple additive composition model improves the state of the art in this task significantly (a gain of nearly 10%). With

weighted additive composition, the gain is more than 2% beating the state-of-the-art in Hindi [30]. When used for the much better-researched, larger datasets of English the system does respectably, but well behind the very best models that attempt more complex composition models. So the question arises as to whether the very significant gains in Hindi are due to some quirk in the dataset, or could it be that Hindi word vectors are particularly informative, e.g. owing to more highly inflected nature of its surface forms. Also, if the results are not corpus-specific, it also raises the possibility that word vector methods may result in significant gains in other similar problems for Hindi.

For example, the sentences below give a brief idea of what positive and negative sentiment means.

- Positive: यह फिल्म अच्छी है
- Negative: यह समान बहुत खराब है

Majority of the existing work in this field is in English (Pang and Lee, 2008). [29] give a summary of work done in Hindi in the field of opinion mining.

We then go a step ahead to build a model using document vectors along with tfidf and ensemble techniques to achieve state-of-the-art result on IMDB movie review dataset achieving a significant improvement over the previous best. We have got an accuracy of 94.19% by using an ensemble method merged with recurrent neural network to show the excellency of generative models merged with discriminative ones. Our model uses concatenation of average word vectors, sentence vectors and tf-idf to build informative document vectors which constitute a feature set for our SVM classifier. We have used skip-gram model (Mikolov et al., 2013b) for building word vectors because deep network model of Collobert et al.(2008) takes too much time for training. Also the word embeddings obtained by them are not as good as those obtained by Mikolov et al.(2013b).

For Hindi, we have experimented with tf-idf and weighted average word vectors for building our feature set. Our model shows slight improvement in performance if

⁰en.wikipedia.org/wiki/List_of_languages_by_number_of_native_speakers

we filter out certain corpus based stop words. This is a first attempt to use word vectors for sentiment analysis in Hindi. Word vectors capture both semantic and syntactic information for a given corpus. Words which are semantically and syntactically related tend to be closer in high-dimensional space.

1.1 The Problem

The problem undertaken in this thesis is to use Distributional Semantics for under-resourced languages such as Hindi to cater to a very significant NLP task of sentiment analysis. We also try to highlight the differences in language and how efficacy of a model changes with language. We also took the task of using distributional semantics and general semantics with various forms of composition of vectors and even features to build good models for English as well. We have also laid a systematic background to the problem of context size and how it is related to syntactic and semantic similarity in text.

1.2 Contributions of this Thesis

There are three main contributions of this thesis:

- We have successfully applied distributional models for Hindi and achieved state-of-the-art results on standard corpus of product and movie reviews. As a result, we have successfully published our work in regICON-2015 organized at IIT-BHU.
- We have also compiled a corpus in Hindi containing around 700 Movie Reviews which is the largest corpus in Hindi in this domain. For reference, we will call it as *700-Movie reviews* dataset.
- We have achieved state-of-the-art results on IMDB Movie Review dataset(94.19%) and have beaten the previous best by over 1.6% in terms of accuracy.

1.3 Organization of this Thesis

The rest of this thesis is organized as follows. Chapter 2 presents related work done and discusses the various background work that one must be acquainted with in order to understand the work presented. Chapter 3 discusses the data sources and presents the statistics of each dataset. Chapter 4 gives an overview of dimensionality reduction techniques and how we have successfully applied them on vectors of our Hindi corpus to achieve very good results. Chapter 5 then discusses in detail the implemented methods and algorithms. Chapter 6 presents a summary of the results that was achieved. Chapter 7 presents concluding remarks to our work and talks about what can be done further.

Chapter 2

Related Work and Background

Sentiment analysis is a well-known research area in NLP today (see reviews in [26] and Pang et al. (2008), and also challenge in SemEval-2014). Early work on movie review sentiments achieved an accuracy of 87.2% (Pang et al. 2004) on a dataset that discarded objective sentences and used text categorization techniques on the subjective sentences. Le and Mikolov (2014) use word vector models and obtain 92.6% accuracy on IMDB movie review dataset. They used distributed bag-of-words model, which they call as *paragraph vector*. More difficult challenges involve short texts with nonstandard vocabularies, as in twitter. Here, some authors focus on building extensive feature sets (e.g. Mohammad et al.(2013); F-score 89.14). Wang et al. (2014) propose a word vector neural-network model, which takes both sentiment and semantic information into account. This word vector expression model learns word semantics and sentiment at the same time as well as fuses unsupervised contextual information and sentence level supervised labels.

Neelakantan et al.(2015)[25] took word vector models to next level where they proposed multiple embeddings per word. Nearly all the work before this assumes a single vector per word type ignoring polysemy and thus jeopardizing their usefulness for downstream tasks. The authors present an extension to the Skip-gram model that efficiently learns multiple embeddings per word type. It differs from recent related work by jointly performing word sense discrimination and embedding learning, by non-parametrically estimating the number of senses per word type, and by its

efficiency and scalability. They achieve state-of-the-art results in the word similarity in context task and demonstrate its scalability by training with one machine on a corpus of nearly 1 billion tokens in less than 6 hours.

Since most sentiment analysis is oriented towards semantics, and one may bypass the syntactic processing which remains poor for Hindi. Methods that have been used are largely based on assigning a sentiment effect for individual words, and then combining these in some manner to come up with an overall sentiment for the document. Such methods ignore word order and have been criticized since the import of a sentence can change completely simply by re-arranging the words, though the sentiment evaluation remains the same. Several groups have attempted to improve the situation by modeling the composition of words into larger contexts [12, 31, 9, 2].

However, most of the work on sentiment analysis in Hindi has not attempted to form richer compositional analyses. For the type of corpora used here, the best results, obtained by combining a sentiment lexicon with hand-crafted rules (e.g. modeling negation and "but" phrases), reach an accuracy of 80% [23].

There has been limited work on sentiment analysis in Hindi – see review in [16], who surveys sentiment analysis in non-English languages). Joshi et al. (2010) compared three approaches: In-language sentiment analysis, Machine Translation and Resource Based Sentiment Analysis. By using WordNet linking, words in English SentiWordNet were replaced by equivalent Hindi words to get H-SWN. The final accuracy achieved by them is 78.1%.

[1] traversed the WordNet ontology to antonyms and synonyms to identify polarity shifts in the word space. Further improvements were achieved by using a partial stemmer (there is no good stemmer / morphological analyzer for Hindi), and focusing on adjective/adverbs (45 + 75 seed words given to the system); their final accuracy was 79.0% for the product review dataset. Mukherjee et al. (2012) presented the inclusion of discourse markers in a bag-of-words model and how it improved the sentiment classification accuracy by 2-4%. Mittal et al. (2013) incor-

porate hand-coded rules dealing with negation and discourse relations and extend the HSWN lexicon with more opinion words. Their algorithm achieves 80.2% accuracy on classification of movie reviews on a separate dataset.

[18] presents new recurrent neural network based language model which has found vital application in various NLP tasks as well. They provide ample empirical evidence to suggest that connectionist language models are superior to standard n-gram techniques, except their high computational (training) complexity. The model has also been extended to speedup training time by few optimizations.(More details in 2.3).

The algorithms and data structures used in this thesis have been introduced and discussed below.

2.1 Word-Vectors and Representations

2.1.1 Skip-Gram

Mikolov et al. (2013b) proposed two neural network models for building word vectors from large unlabeled corpora; Continuous Bag of Words(CBOW) and Skip-Gram. In the CBOW model, the context is the input, and one tries to learn a vector for the central word; in Skip grams, the input is the target word and one tries to guess the set of contexts. The Skip gram was found to perform better on smaller corpora, and here we have focused on this model for building our word vectors. The model uses each current word as an input to a log-linear classifier with continuous projection layer, and predict words within a certain range before and after the current word. The objective is to maximize the probability of the context given a word within a language model:

$$p(c|w; \theta) = \frac{\exp^{v_c \cdot v_w}}{\sum_{c' \in C} \exp^{v_{c'} \cdot v_w}}$$

where v_c and $v_w \in R^d$ are vector representations for context c and word w respectively. C is the set of all available contexts. The parameters θ are $v_{c,i}$, $v_{w,i}$ for $w \in V$,

$c \in C, i \in 1, \dots, d$ (a total of $|C| \times |V| \times d$ parameters).

The weights between the input layer and the output layer can be represented by a $V \times N$ matrix \mathbf{W} . Each row of \mathbf{W} is the N -dimension vector representation v_w of the associated word of the input layer. Given a word, assuming $x_k = 1$ and $x_{k'} = 0$ for $k' \neq k$, then

$$h = x^T W = W_{(k, \cdot)} := v_{w_I}$$

which is essentially copying the k -th row of \mathbf{W} to \mathbf{h} . v_{w_I} is the vector representation of the input word w_I .

From the hidden layer to the output layer, there is a different weight matrix $\mathbf{W}' = \{w'_{ij}\}$ which is a $N \times V$ matrix. Then we can use soft-max, a log-linear classification model, to obtain the posterior distribution of words, which is a multinomial distribution.

On the output layer, instead of outputting one multinomial distribution, we are outputting C multinomial distributions. Each output is computed using the same hidden→output matrix.

$$p(w_{c,j} = w_{O,c} | w_I) = y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})}$$

where $w_{c,j}$ is the j -th word on the c -th panel of the output layer; $w_{O,c}$ is the actual c -th word in the output context words; w_I is the only input word; $y_{c,j}$ is the output of the j -th node on the c -th panel of the output layer; $u_{c,j}$ is the net input of the j -th node on the c -th panel of the output layer. Because the output layer panels share the same weights, thus

$$u_{c,j} = u_j = v_{w_j}^T \cdot h, \text{ for } c = 1, 2, \dots, C$$

where v_{w_j}' is the output vector of the j -th word in the vocabulary, w_j , and also v_{w_j}' is taken from a column of the hidden→output weight matrix, \mathbf{W}' .

The loss function is

$$E = -\log p(w_{O,1}, w_{O,2}, \dots, w_{O,C} | w_I)$$

$$\begin{aligned}
&= -\log \prod_{c=1}^C \frac{\exp(u_{c,j_c^*})}{\sum_{j'=1}^V \exp(u_{j'})} \\
&= -\sum_{c=1}^C u_{j_c^*} + C \cdot \log \sum_{j'=1}^V \exp(u_{j'})
\end{aligned}$$

where j_c^* is the index of the actual c -th output context word in the vocabulary.

After taking the necessary derivatives, the update equation for the hidden→output weight matrix, \mathbf{W}' ,

$$w'_{ij}^{(new)} = w'_{ij}^{(old)} - \eta \cdot EI_j \cdot h_i$$

or,

$$v'w_j^{(new)} = v'w_j^{(old)} - \eta \cdot EI_j \cdot \mathbf{h} \text{ for } j = 1, 2, \dots, V$$

The update equation for the input→hidden weight matrix, \mathbf{W} ,

$$vw_I^{(new)} = vw_I^{(old)} - \eta \cdot EH$$

where EH is a N -dimensional vector. Its each component is defined as

$$EH_i = \sum_{j=1}^V EI_j \cdot w'_{ij}$$




Figure 2.1: Skip Gram Model(Figure from Rong (2014))

The main advantage of using skip-gram is that it is computationally less expensive than other neural language models with a complexity of $O(\log V)$ instead of $O(V)$. They use hierarchical softmax [24] to achieve such computational efficiency.

2.1.2 tf-idf

Let $D = d_1, d_2, d_3, \dots, d_N$ be N documents under study and $T = t_1, t_2, t_3, \dots, t_M$ be the M unique terms in these documents, then each document can be represented as a M -dimensional vector:

$$t_d = \{tf_1, tf_2, tf_3, \dots, tf_M\}$$

$tf - idf$ weights discounts the frequency of terms based on their importance to a particular document in the entire document set collection under consideration. This is done as follows:

$$tfidf(d, t) = tf(d, t) \times \log\left(\frac{|D|}{df(t)}\right)$$

Here $df(t)$ is the number of documents in which term t appears.

2.1.3 Vector Averaging for phrases

We will describe the process for Hindi but the same flow was followed for English as well. As an output of the word vector learning, we now have a n -dimensional vector representation for each word in the Hindi corpus. Now we need to assign features for sentences and paragraphs taken from the sentiment dataset (training and test). Mikolov et al. (2013b) and Levy et al. (2014) show that many relational similarities can be recovered by means of vector arithmetic in the embedded space. Thus, additive models are useful, though others have claimed that multiplicative models correlate better with human judgments [22, 31]. In this work, we have retained the simplicity of vector averaging to model larger chunks of discourse. This models the sentence/document in the same high dimensional space.

A preprocessing step involved removing some words that appear at very high or very low frequencies in the corpus. Our model was trained on the Hindi Wikipedia

dump to create vector representations for words. The previous two vectors were concatenated to create another feature set for training purpose.

Algorithm

1. Input the Hindi text corpus
2. Train skip-gram model to obtain word vector representation
3. Given a sentiment training set, obtain average vector data for each sentence/document
4. Obtain tf-idf vector for each sentence/document in the corpus
5. Concatenate vectors of step 3 and step 4 to obtain a feature set for a training instance
6. Train linear SVM with m -fold cross validation to create a classifier (here $m=20$)

2.2 Document Vectors

This distributed representation of sentences and documents [12] modifies word2vec (Skip-Gram) algorithm to unsupervised learning of continuous representations for larger blocks of text, such as sentences, paragraphs or entire documents. The algorithm represents each document by a dense vector which is later trained and tuned to predict words in the sentencedocument.

In Paragraph Vector framework, every paragraph is mapped to a unique vector and id, represented by a matrix D , which is a column matrix. Every word is mapped to a unique vector and word vectors are concatenated or averaged to predict the context, i.e., the next word.

The change in this framework is that the h (in Skip-Gram model's equation) is now constructed in a different way. It is now constructed using both W and D .

The contexts are fixed-length and sampled from a sliding window over the paragraph. The paragraph vector is shared across all contexts generated from the same

paragraph but not across paragraphs. The word vector matrix W , however, is shared across paragraphs. i.e., the vector for "good" is the same for all paragraphs.

The paragraph vectors and word vectors are trained using stochastic gradient descent and the gradient is obtained via backpropagation. At every step of stochastic gradient descent, one can sample a fixed-length context from a random paragraph, compute the error gradient from the network in Figure 2.3 and use the gradient to update the parameters in our model. At prediction time, one needs to perform an inference step to compute the paragraph vector for a new paragraph. This is also obtained by gradient descent. In this step, the parameters for the rest of the model, the word vectors W and the softmax weights, are fixed.




Figure 2.2: Framework for learning word vectors(Figure from Le (2014)).

In Figure 2.2, context of three words ("the", "cat" and "sat") is used to predict the fourth word ("on"). The input words are mapped to columns of the matrix W to predict the output word(Figure from Le (2014)).

In Figure 2.3, the only change is the additional paragraph token that is mapped to a vector via matrix D . In this model, the concatenation or average of this vector with a context of three words is used to predict the fourth word. The paragraph vector represents the missing information from the current context and can act as a memory of the topic of the paragraph.




Figure 2.3: Framework for learning paragraph vectors(Figure from Le (2014)).

The advantage of using paragraph vectors is that they inherit the property of word vectors, i.e., the semantics of the words. In addition, they also take into consideration a small context around each word which is in close resemblance to the n-gram model with a large n . This property is crucial because the n-gram model preserves a lot of information of the sentence/paragraph, which includes the word order also. This model also performs better than the Bag-of-Words model which would create a very high-dimensional representation that has very poor generalization.

2.3 Recurrent Neural Network

The primary feature of a Recurrent Network is that it contains atleast one feed-back connection which captures dynamic temporal behavior and allows it to learn sequences. They have application in tasks such as prediction of a word given the context, perform sequence recognition/reproduction.

Recurrent Neural Networks have many different forms. One of them is a Fully Recurrent Network, a network of neuron-like units, each with a directed connection to every other unit.

The architecture of a simple RNN consists of an input layer x , a hidden layer




Figure 2.4: Simple Recurrent Neural Network(Figure from [18]).

s (also known as context layer or network state) and an output layer y . Since the training is time dependent, we will denote input x as $x(t)$, output y as $y(t)$ and state s as $s(t)$. Input vector $x(t)$ is a concatenation of vector w representing current word, and output from neurons in context layer s at time $t - 1$ (We can also include output from context layer before $t - 1$ as well). The equations of each layer are described below:

$$x(t) = w(t) + s(t - 1) \quad (2.1)$$

$$s_j(t) = f\left(\sum_i x_i(t)u_{ji}\right) \quad (2.2)$$

$$y_k(t) = g\left(\sum_j s_j(t)v_{kj}\right) \quad (2.3)$$

where $f(z)$ is sigmoid function for activation and $g(z)$ is a softmax function for output prediction:

$$f(z) = \frac{1}{1 + e^{-z}} \quad (2.4)$$

$$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}} \quad (2.5)$$

[18] claim that size of hidden layer reflects amount of training data with smaller size leading to less number of layers. Weights are initialized to random small values

and updated using gradient descent. Output layer $y(t)$ represents probability distribution of next word given previous word $w(t)$ and context $s(t-1)$. The objective function is:

$$error(t) = actual(t) - y(t) \quad (2.6)$$

where $actual(t)$ is a 1-hot vector representing the word that should have been predicted given the context. Table 2.1 represents results of [18] obtained in WSJ ex-

Model	DEV WER	EVAL WER
Lattice 1 best	12.9	18.4
Baseline-KN5 (37M)	12.2	17.2
Discriminative LM (37M)	11.5	16.9
Joint LM (70M) (37M)	-	16.7
Static 3xRNN + KN5 (37M)	11.0	15.5
Dynamic 3xRNN + KN5 (37M)	10.7	16.3

Table 2.1: Comparison of WSJ results obtained with various models(RNN is trained on just 6.4M words)

periments using RNN.

[19] present several modifications of the original recurrent neural network language model (RNN LM). The present approaches that lead to more than 15 times speedup for both training and testing phases. They also show the importance of backpropagation through time(BPTT) which is an extension to backpropagation algorithm for recurrent networks. With truncated BPTT, the error is propagated through recurrent connections back in time for a specific number of time steps. Thus, the network learns to remember information for several time steps in the hidden layer when it is learned by the BPTT.

The speedup is obtained by assuming that words can be mapped to classes. Thus if we assume that each word belongs to exactly one class, we can first estimate the probability distribution over each class using RNN and then calculate the probability of a particular word from the desired class assuming unigram distribution of words within the class. Thus, now we are reducing the connections between hidden

and output layer from V to C which is a significant improvement.

They are often very sensitive to small changes in its parameters which changes the gradient by a large amount. Few others are, Bi-directional RNN and Hierarchical RNN.

2.4 Semantic Composition

The Principle of Compositionality is that meaning of a complex expression is determined by the meaning of its parts or constituents and the rules which guide this combination. It is also known as *Frege's Principle*. In our case, the constituents are word vectors and the expression in hand is the sentence/document vector. For example,

The movie is funny and the screenplay is good

In the above sentence, consider the word vectors are represented by $w(x)$ and the sentence vector as $S(x)$. Hence,

$$S(x) = c_1 w_1(x) \Theta c_2 w_2(x) \Theta c_3 w_3(x) \Theta c_4 w_4(x) \dots \Theta c_k w_k(x) \quad (2.7)$$

where Θ can be any operation (e.g., addition, multiplication) and c_i s are constants.

Composition	Accuracy
Multiplication	50.30
Average	88.42
Weighted Average	89.56

Table 2.2: Results of Vector Composition with different Operations

Analyzing the results from Table 2.2, we observed that when we deal with large number of features, there is a presence of large number of *zeros* and presence of a single zero in a feature will make that features contribution zero in the final vector, which happens in our case and thus multiplicative composition fails.

We, therefore, adopt both simple and weighted average methods in our work. The

advantage with addition is that, it doesnot increase the dimension of the vector and captures high level semantics with ease. In fact, [37] have used simple average to construct phrase vectors which they have later used to find phrase level similarity using cosine distance.

[21] showed that relations between words are reflected to a large extent in the offsets between their vector embeddings. They also use additive composition to reflect semantic dependencies.

$$\text{queen} - \text{king} \approx \text{woman} - \text{man}$$

[3] clearly show that vectors of Neural Language Model and Distributed Model when used with additive composition outperform those with multiplicative composition in Paraphrase Classification task. DM vectors outperform by nearly giving accuracy difference of 6%. They also perform very well on Phrase similarity tasks.

[31] also present yet another model for semantic composition but that uses a sentiment treebank which is a very costly structure to build and it is task dependent. For under-resourced languages such as Hindi, it would take years to build (for English, task done through Amazon Mechanical Turk). This leads to appreciation of models such as Additive Composition.

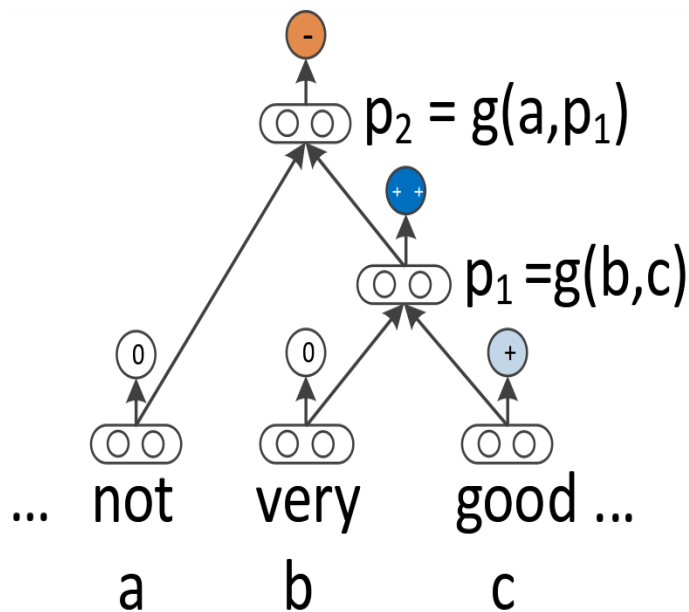


Figure 2.5: Approach of Recursive Neural Network(Figure from [31]).

Figure 2.5 depicts the approach of recursive neural network. When an n -gram is given to the compositional models, it is parsed into a binary tree and each leaf node, corresponding to a word, is represented as a vector. Recursive neural models will then compute parent vectors in a bottom up fashion using different types of compositionality functions g . The parent vectors are again given as features to a classifier.

Chapter 3

Data Acquisition

This section describes the corpus used in our experiments.

3.1 Hindi Corpus

3.1.1 Product and Movie Review Corpus

We experimented on two Hindi review datasets. One is the Product Review dataset (LTG, IIIT Hyderabad) containing 350 Positive reviews and 350 Negative reviews. The other is a Movie Review dataset (CFILT, IIT Bombay) containing 127 Positive reviews and 125 Negative reviews.

Each review is around 1-2 sentences long and the sentences are mainly focused on sentiment, either positive or negative.

3.1.2 700-Movie Review Corpus(Our Contribution)

We collected Hindi movie reviews from websites such as *Dainik Jagran* and *Navbharat Times*. The movie reviews are longer than the previous corpus and contains subjects other than sentiment. There are in total 697 movie reviews from both the websites. The statistics compiled is described below.

Dainik Jagran	
Positive Reviews	210
Negative Reviews	226
Total Reviews	436
29.7 sentences per document	
427.1 words per document	

Table 3.1: Statistics of Dainik Jagran Movie Reviews

Navbharat Times	
Positive Reviews	146
Negative Reviews	115
Total Reviews	261
29.7 sentences per document	
607.2 words per document	

Table 3.2: Statistics of Navbharat Times Movie Reviews

Overall	
Positive Reviews	356
Negative Reviews	341
Total Reviews	697
29.7 sentences per document	
494.6 words per document	

Table 3.3: Statistics of Movie Reviews from Jagran and Navbharat

3.1.3 Wikipedia

We also trained our skip-gram model on Hindi Wikipedia text dump (approx. 290MB) containing around 24M words with 724K words in the vocabulary. This provided us with good embeddings due to larger size and contents from almost all domains.

3.2 English Corpus

3.2.1 IMDB Movie Review

We trained on IMDB movie review dataset (Maas et al.(2013)) which consists of 25,000 positive and 25,000 negative reviews. The core dataset contains 50,000 reviews split evenly into 25k train and 25k test sets. The overall distribution of labels is balanced (25k pos and 25k neg). It also contains an additional 50,000 unlabeled documents for unsupervised learning.

In the entire collection, no more than 30 reviews are allowed for any given movie because reviews for the same movie tend to have correlated ratings. Further, the train and test sets contain a disjoint set of movies, so no significant performance is obtained by memorizing movie-unique terms and their associated with observed labels. In the labeled train/test sets, a negative review has a score ≤ 4 out of 10, and a positive review has a score ≥ 7 out of 10. Thus reviews with more neutral ratings are not included in the train/test sets. In the unsupervised set, reviews of any rating are included and there are an even number of reviews > 5 and ≤ 5 . For example,

- **Positive:** After reading previews for this movie I thought it would be a let down, however after I got my region 1 dvd (the dvd was available before the film hit the uk cinemas) I was pleasantly surprised, strong performances from all cast members make this a very enjoyable movie. The fact that the script is quite weak means that you dont get bogged down in story and therefore the repeat viewing factor is greater. I recommend this movie to one and all
- **Negative:** This is by far the worst non-English horror movie I've ever seen. The acting is wooden, the dialogues are simply stupid and the story is totally braindead. It's not even scary. 2 out of 10 from me.

3.2.2 Trip Advisor Review

It contains around 240K reviews(206MB) from hotel domain. Reviews with overall rating ≥ 3 were annotated as positive and those with overall rating < 3 were annotated as negative.

There were total 27315 words in the vocabulary after removing those with count < 10 . Overall, there were 154448 words in the corpus.

The dataset was split into 80-20 ratio for training and testing purpose.

The Meta data includes: Author, Content, Date, Number of Reader, Number of Helpful Judgment, Overall rating, Value aspect rating, Rooms aspect rating, Location aspect rating, Cleanliness aspect rating, Check in/front desk aspect rating, Service aspect rating and Business Service aspect rating. Ratings ranges from 0 to 5 stars, and -1 indicates this aspect rating is missing in the original html file.

3.2.3 Amazon Review[15]

Reviews with overall rating ≥ 3 were annotated as positive and those with overall rating < 3 were annotated as negative. The dataset was split into 80-20 ratio for training and testing purpose.

There were 3 review datasets: Watches, Electronics and MP3 each of size 30.8MB, 728.4MB and 27.7MB respectively.

Electronics dataset consists of 1,241,778 reviews, Watches Dataset consists of 68,356 reviews and MP3 Dataset consists of 31,000 reviews.

3.2.4 Wikipedia

We also trained our skip-gram model on English Wikipedia text dump (approx. 20.3GB) containing around 3.5B words with 7.8M words in the vocabulary. This provided us with good embeddings due to larger size and contents from almost all domains.

Chapter 4

Dimensionality Reduction

Dimensionality Reduction is the process of reducing the number of random variables in such a way that the remaining variables effectively reproduce most of the variability of the dataset. The reason for using such techniques is because of the *curse of dimensionality* which is a phenomena that occurs in high-dimension but doesn't occur in low-dimension.

In machine learning problems that involve learning a "state-of-nature" (maybe an infinite distribution) from a finite number of data samples in a high-dimensional feature space with each feature having a number of possible values, an enormous amount of training data are required to ensure that there are several samples with each combination of values. With a fixed number of training samples, the predictive power reduces as the dimensionality increases, and this is known as the Hughes effect[3] or Hughes phenomenon (named after Gordon F. Hughes)¹

It happens, in general, that the intrinsic dimension is small but the data is represented by a large number of features. e.g. there may be some features across which the variance is very low and hence we can safely ignore such features reducing the dimension.

Dimensionality Reduction helps in visualization of high-dimensional data in 2D or 3D, data compression for efficient retrieval and storage and noise removal. In machine learning, it is often used for feature selection and feature extraction. In this

¹Wikipedia

work, we have successfully used dimensionality reduction for feature selection i.e., we have selected an optimal subset of features from the given data to achieve a large improvement in accuracy.

Next we will discuss two of the techniques which we have used in our work and present the results obtained after applying these techniques.

4.1 Principal Component Analysis(PCA)

Principal Component Analysis(PCA) is a very famous technique for dimensionality reduction in machine learning. When given a data with n -dimensions/features, PCA aims to find a linear subspace of dimension d lower than n such that this reduced space contains most of the data points and it maintains most of the variability of the data.

It can be easily proved that first principal component is given by the normalized eigenvector with the largest associated eigenvalue of the sample covariance matrix S . A similar argument can show that the d dominant eigenvectors of covariance matrix S determine the first d principal components. Also, the projection onto the principal subspace minimizes the squared reconstruction error, $\sum_{i=1}^t \|x_i - \hat{x}_i\|^2$.

Algorithm 1: PCA Algorithm

- 1 **Recover basis:** Calculate $XX^T = \sum_{i=1}^t x_i x_i^T$ and let $U =$ eigenvectors of XX^T corresponding to the top d eigenvalues.
- 2 **Encode training data:** $Y = U^T X$ where Y is a $d \times t$ matrix of encodings of the original data.
- 3 **Reconstruct training data:** $\hat{X} = UY = UU^T X$.
- 4 **Encode test example:** $y = U^T x$ where y is a d -dimensional encoding of x .
- 5 **Reconstruct test example:** $\hat{x} = Uy = UU^T x$.




Figure 4.1: PCA of a multivariate Gaussian distribution centered at $(1, 3)$ with a standard deviation of 3 in roughly the $(0.878, 0.478)$ direction and of 1 in the orthogonal direction. The vectors shown are the eigenvectors of the covariance matrix scaled by the square root of the corresponding eigenvalue, and shifted so their tails are at the mean (Figure from [36]).

4.2 ANOVA: F-value

Analysis of variance (ANOVA) is a set of statistical models used to analyze the differences between the group means and variation among and between groups. It was developed by R.A. Fisher. It is also a measure of statistical significance of the data point with reference to other data points present in the sample.

In our experiment, we have used F-test to reduce the number of features and it has led to a considerable increase in classification accuracy on our new movie dataset (Refer 3.1.2).

The F-statistic is calculated as:

$$F = \frac{\text{variance between groups}}{\text{variance within groups}}$$

The variance between groups is given as:

$$\frac{\sum_i n_i (\bar{Y}_i - \bar{Y})^2}{K - 1} \quad (4.1)$$

where \bar{Y}_i denotes the sample mean in the i^{th} group, n_i is the number of observations in the i^{th} group, \bar{Y} denotes the overall mean of the data, and K denotes the number of groups.

The variance within groups is given as:

$$\frac{\sum_{ij} (Y_{ij} - \bar{Y}_i)^2}{N - K} \quad (4.2)$$

where Y_{ij} is the j^{th} observation in the i^{th} out of K groups and N is the overall sample size.

4.3 Result

Method	Feature Selection	Accuracy
Document Vector	None	74.57
	PCA(n=50)	76.33
	ANOVA-F	88.07
Weighted Average Word Vector	None	76.43
	ANOVA-F	90.37
	PCA(n=50)	78.61

Table 4.1: Accuracies on 700-Movie Review Dataset

Table 4.1 summarizes how feature selection has improved classification accuracy on the 700 Movie review dataset. With ANOVA-F, we selected around 4k features but with PCA, this number was just 50. So, the low accuracy with PCA can be attributed to the fact that we may have lost some important features in low dimen-

sion. Also, PCA cannot work with size of dimension $d > \text{size of learning set}$.

This sharp decrease in accuracy in both cases: Document vector and Weighted Average word vector happens because ANOVA-F selects features with larger variance across group and thus reduces noise to a larger extent whereas PCA reduces angular variance which is not effective in this case due to the distribution of data points in high-dimensional space.

Corpus	Feature Selection	Accuracy
Movie Reviews-IITB	None	79.17
	ANOVA-F	89.17
Product Reviews-IIIT	None	86.86
	ANOVA-F	90.83

Table 4.2: Accuracies on IITB Movie Review and IIIT Product Review Dataset using Document Vectors

Again from Table 4.2, we observe that ANOVA-F performs much better than PCA. Reducing dimension to 4k preserves important features whereas we are losing information once that number is reduced to 50.

So we easily infer that with these datasets, ANOVA-F is the feature selection method to adopt.

Chapter 5

Models and Experiments

5.1 Context Size

[20] describe in their work how context size affects the type and quality of word vectors. One of the observations is that a smaller context size tend to capture syntactic similarity to a much better extent. As we increase the context size, vectors tend to capture more of semantic similarity.

Figure 5.1 shows how the high dimensional representation changes when we change the context size. The embeddings were formed with context sizes 5 and 10 by training on Latest Wikipedia dump. We see that 5.1(a), has words $\{high, higher, highest\}$ very close to each other which is in accordance to our argument that smaller context size promotes syntactical similarity whereas these words are farther when we increase the context size. From 5.1(b), we can see that there is a larger semantic difference between two clusters of words: $\{wiki\ project, wikipedia\}$ and $\{facebook, html, id\}$. This again justifies our argument that larger context size favors more of semantic similarity than syntactic similarity.

We try to look how context size affects the accuracy on two amazon datasets, Watches and MP3. Figure 5.2 depicts the accuracies obtained by varying the context size from 5 to 10. In both the cases, we find a peak at context size 7. There is an increase from 5 to 7 which denotes increase of context strength. As a result, we see increase in accuracies. Once the context size goes beyond a threshold, the




Figure 5.1: High Dimensional representation of Wiki Text with a) Context Size 5
b) Context Size 10




Figure 5.2: Variation of Accuracy with Different Context Size on Watches and MP3 Datasets.

semantic counterpart dominates which attains maximum at context size 10.

We also see that larger data size or training data affects the accuracy on test set.

In this case, *Watches* dataset is larger than *MP3* dataset.

5.2 SkipGram or CBOW

SkipGram model tends to predict a context given a word whereas CBOW model predicts a word given a context. It seems intuitive and also from observation [17] that SkipGram will perform better on semantic tasks and CBOW on syntactic tasks. We now try to evaluate how they differ on classification accuracies on the two datasets: *Watches* and *MP3*. Figure 5.3 show that skipgram outperforms CBOW on sentiment classification task. It can be justified by the fact that sentiment inclination of a document is more oriented towards semantics of that document rather than just




Figure 5.3: Variation of Accuracy with skipgram and cbow on Watches and MP3 Datasets.

syntax and our results clearly demonstrate this fact.

5.3 Skip-Gram and *tf-idf* based Word Vectors

In this experiment, we first generated n -dimensional word vectors by training skip-gram model on product and movie review corpus in Hindi and IMDB corpus in English. We then averaged-out word vectors for each document to create document vectors. This now acts as a feature set for that particular document. We also created *tf-idf* vectors for each document. This can be seen as a vector representation of that particular document. We then concatenated these document vectors with document vectors obtained after averaging-out word vector of each document.

In another experiment, we filtered out stop-words on the basis of their frequency in the corpus. Words which had very high or very low frequency were pruned as they had negligible contribution to the sentiment polarity of a document. This is a

noise-reduction step and gives better results.

5.4 Weighted Average Word Vectors

Algorithm 2 presents a new approach to build document vectors using distributional semantics. Till date, everyone has ignored how to effectively use vector composition techniques and as a result, this area has seen very less attention. But we have successfully used *idf* values to give weights to word vectors and hence obtain much better sentence/document vectors. Refer to 6.2 to see how this new technique has improved sentiment classification on Hindi reviews.

Algorithm 2: Weighted Average for Vector Composition

- 1 **Skip-Gram:** Train skip-gram model on text to obtain word vectors.
- 2 **idf Calculation:** Calculate *idf* value of each word with respect to given document (Refer 2.1.2).
- 3 **Additive Composition:** For each document,

$$\sum_i (idf_i \times WordVector_i)$$
 is the document vector.
- 4 **Feature Set:** Repeat 3 for each document in training and test set to build document vectors.
- 5 **Classification:** Use a binary classification algorithm (SVM or Logistic Regression) to classify each test sample as positive or negative.

The advantage of this model is that once we obtain *idf* values from training corpus, we can directly use it with test corpus without any additional computation.

5.4.1 Graded Weighting

We describe two approaches to incorporate graded weighting into word vectors for building document vectors. Let v_{w_i} be the vector representation of the i^{th} word. Then document vector v_{d_i} for i^{th} document is:

$$v_{d_i} = \begin{cases} 0 & w_k \in stopwords \\ \sum_{w_k \in d_i} v_{w_k} & w_k \notin stopwords \end{cases}$$

The above equation is 0-1 step-function which ignores contribution of all stop words. Now we propose another schema which weighs the contribution of each word while building document vector with a graded approach. We define $idf(t, d) = \log(\frac{|D|}{df(t)})$ where t is the term, d is the document and other notations are same as in previous subsection. The new document vector representation considering this graded schema is:

$$v_{d_i} = \begin{cases} 0 & idf(w_k, d_i) \leq \delta \\ \sum_{w_k \in d_i} idf(w_k, d_i) \cdot v_{w_k} & otherwise \end{cases}$$

where δ is a pre-defined threshold below which the word has no importance and above which the idf terms gives importance to that particular word.

Till date, everyone has ignored how to effectively use vector composition techniques and as a result, this area has seen very less attention. But we have successfully used idf values to give weights to word vectors and hence obtain much better sentence/document vectors. The advantage of this model is that once we obtain idf values from training corpus, we can directly use it with test corpus without any additional computation.

Tables 5.1 demonstrate the effectiveness of our proposed graded weighting technique. Without $tf-idf$ features, our proposed method performs better in the graded idf weighting case and when we include $tf-idf$ features, 0-1 weighting perform better than idf graded technique and both perform better than the previous state-of-the-art. We see that with larger weights there is a decrease in accuracy and that is

because we are now filtering out more words which are important while building document vector.

Method	Weight	Accuracy(1)	Accuracy(2)
0-1 Weighting	0	93.84	93.06
	1	93.91	93.18
Graded idf Weighting	2	93.89	93.17
	2.5	93.87	93.16
	2.8	93.86	93.16
	3	93.86	93.22
	4	93.83	93.12
	5	93.75	93.03

Table 5.1: Results on IMDB Movie Reviews using Various Weighting Techniques(Composite Document Vector);Accuracy(2) is when we exclude tf-idf features

5.5 Word and Document Vectors with tf-idf: Enhanced Document Vectors

This experiment redefined document representation in NLP used for sentiment classification. It has the property of including both syntactic and semantic properties of a piece of text. The limitations of skip-gram word vectors have been fulfilled by document vectors and hence we achieve state-of-the-art results on IMDB movie review dataset.

5.6 Ensemble of RNNLM and Enhanced Document Vectors

This experiment was done on IMDB movie review dataset. Here, we first trained a *recurrent neural network* and then obtained predictions on test reviews in terms of probability. We trained another classifier(Linear SVM) using Enhanced Document Vectors and then obtained predictions on test reviews. We then merged these two predictions using a simple heuristic, described below, to obtain final classification. This led to a tremendous increase in classification accuracy(Refer section 6).

Let y^* be actual output and y is the predicted output. The heuristic used is:

$$((RNNLM_{pred} - 1) * 7 + (0.5 - SVM_{pred})).y < 0 \text{ if } y^* = 1 \quad (5.1)$$

$$((RNNLM_{pred} - 1) * 7 + (0.5 - SVM_{pred})).y < 0 \text{ otherwise} \quad (5.2)$$

5.7 Embeddings

The quality of word vectors can be evaluated by comparing them with words which are closer to them semantically and syntactically. This is usually done via cosine similarity.

$$similarity = \cos(\theta) = \frac{A.B}{\|A\|\|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}} \quad (5.3)$$




Figure 5.4: t-SNE visualization of the top 5000 Hindi Words in high dimensional space. (Magnify to see details).

Another evaluation can be done through tSNE [34] which helps in visualization

which maps each high-dimensional data point to a two or three-dimensional map. In our experiment, we took 5K words and plotted them with tSNE (fig. 5.4).


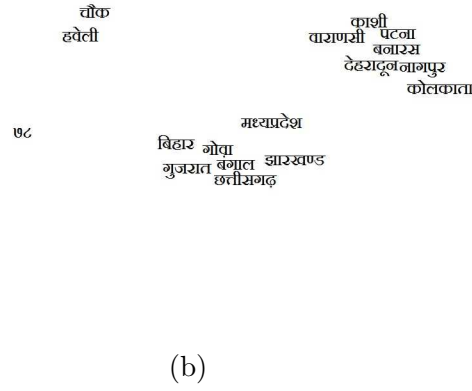



Figure 5.5: A closer look at two clusters in the visualization showing a) quantity relations, b) locations and c) diseases.

Figure 5.5 gives a closer look into few clusters which depicts the relation between words in high dimensional space. Figure 5.5a shows that words such as मौजूद and

उपलब्ध are closer to each other but farther from words such as ज्यादा and अधिक.

5.8 Work Flow

This section summarizes our approach in brief with graphics. This is a general classification system and includes building of enhanced document vectors.

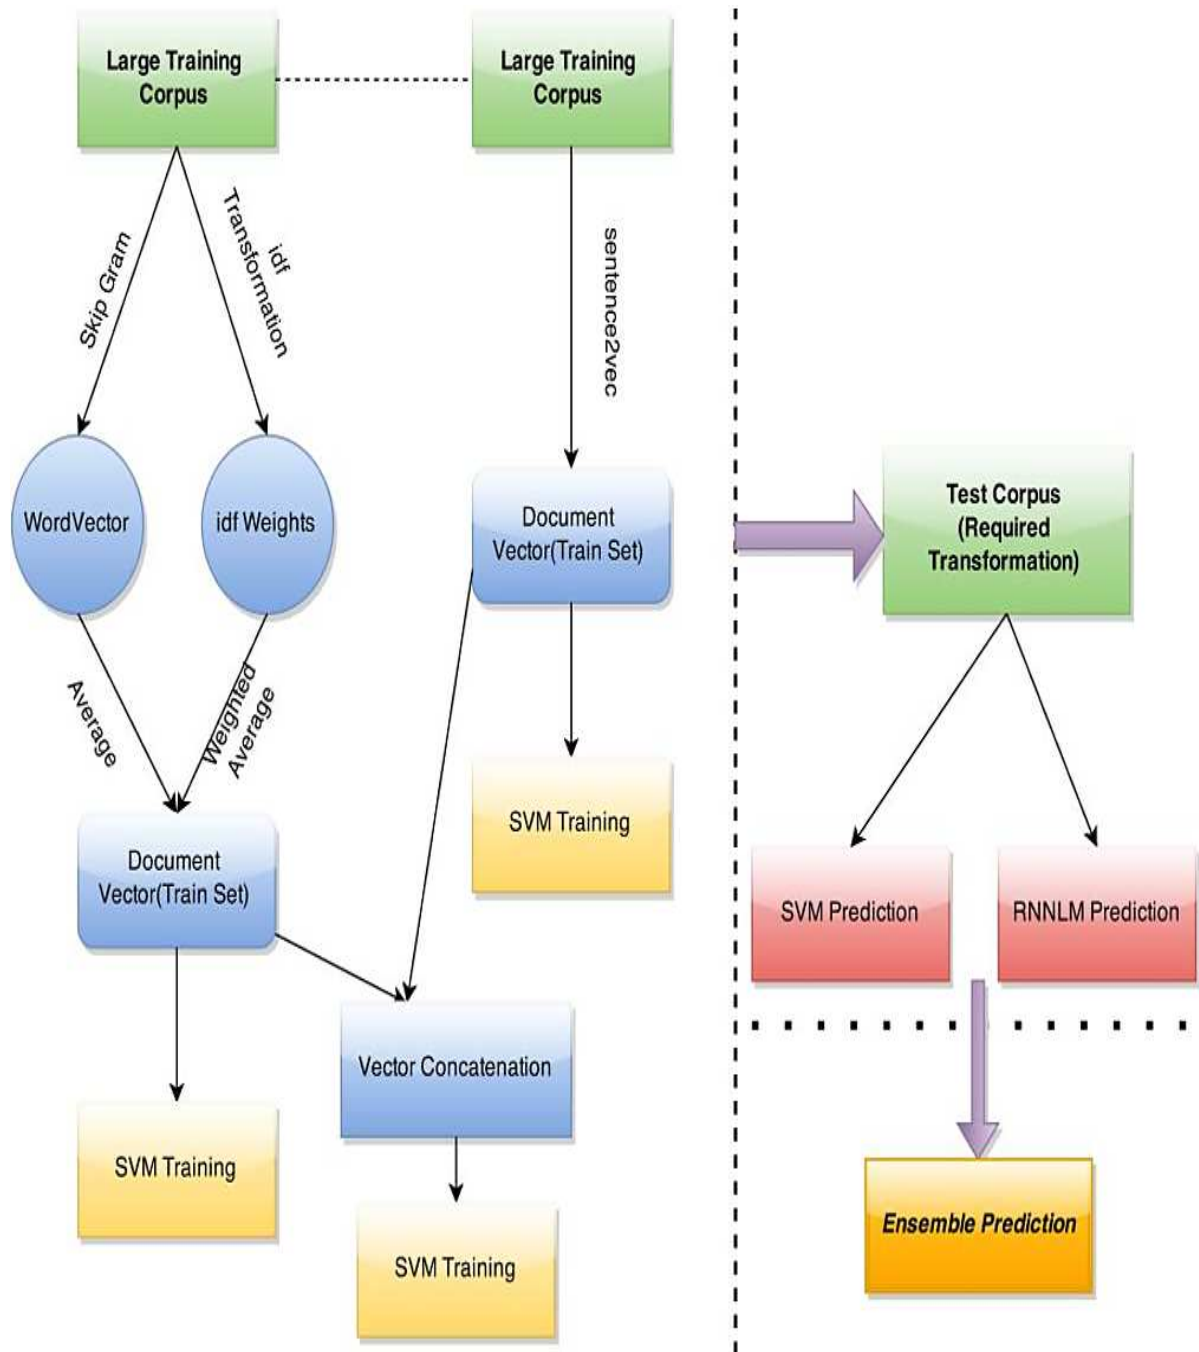


Figure 5.6: Work Flow Summary.

5.9 Technical Specifications

5.9.1 Tools/Libraries

- Eclipse: An Integrated Development Environment helped speed up the process of coding and its subsequent debugging.
- Python: Being a very common and widely used programming language, loads of documentation and third party libraries are available.
- Scikit: Open source library built on NumPy, SciPy and matplotlib which contains simple and efficient tools for data mining and data analysis.
- Gensim: It is an open-source vector space modeling and topic modeling toolkit, implemented in the Python programming language, using NumPy, SciPy and optionally Cython for performance. It is specifically intended for handling large text collections, using efficient online algorithms.

5.9.2 System Requirements

The code and tests have been successfully run on the following configuration. Any system with a configuration equal or higher than this should be able to do the job faster.

- RAM: 16GB; fails to run on 8GB, due to the high amount of in memory data.
- CPU: Intel Core i7, 4th Generation; have used all 8 virtual cores with hyper-threading, with lower CPU, processes would take longer to complete.
- HDD: 1TB dedicated; Wikipedia Dump is the only component using considerable persistent memory. Storing pickle dumps of various datasets consumes considerable amount of memory.

Chapter 6

Results

6.1 Results on English Datasets

Method	IMDB	Amazon	Hindi
RNNLM (Baseline)	86.45	90.03	78.84
Paragraph Vector [12]	92.58	91.30	74.57
Averaged Vector	88.42	88.52	79.62
Weighted Average Vector	89.56	88.63	85.90
Composite Document Vector	93.91	92.17	90.30

Table 6.1: Comparison of accuracies on 3 Datasets(IMDB, Amazon Electronics Review and Hindi Movie Reviews(IITB)) for various types of document composition models. The state of the art for these tasks are: IMDB: 92.58% [12]; Amazon:85.90% [6], Hindi:79.0% [1].

Table 6.1 gives an overview of how our method compares against the earlier methods in sentiment analysis on 3 different datasets. Even the word vector averaging seems to perform quite well when compared against complex baselines such as RNNLM.

Method	Accuracy
Maas et al.(2011)	88.89
NBSVM-bi (Wang & Manning, 2012)	91.22
NBSVM-uni (Wang & Manning, 2012)	88.29
SVM-uni (Wang & Manning, 2012)	89.16
Paragraph Vector (Le and Mikolov(2014))	92.58
WordVector+Wiki(Our Method)	88.60
Weighted WordVector+TfIdf(Our Method)	89.56
Weighted WordVector+TfIdf+Document Vector	93.91

Table 6.2: Results on IMDB Movie Review Dataset

Table 6.2 summarizes the results obtained by others and by us on the IMDB movie review dataset. We have gone above the previous best [12] by a margin of 1.33% using enhanced document vector.

Method	Accuracy
WordVector Averaging+TfIdf+Document Vector	93.91
WordVector+TfIdf+Document Vector+RNNLM(Our Method)	94.19

Table 6.3: Results on IMDB Movie Review Dataset

The main contributor for improvement in results is our enhanced document vector which overcomes the weaknesses of BOW and document vectors taken separately. Table 6.3 is a further improvement in results once we incorporate predictions of RNNLM and enhanced document vector model together(voting ensemble).

Table 6.4 gives an overview of the results obtained with various experiments that we conducted on IMDB dataset. We have tried various vector composition models and feature composition models to obtain good results. Even a simple feature composition of word and document vector yield an improvement over the existing state-of-the-art suggesting that feature composition is a vital area when it comes to NLP tasks such as sentiment analysis.

Method	Accuracy
WordVector Averaging	88.42
Weighted WordVector Average	89.56
Weighted WordVector Averaging+Wiki	88.60
Weighted WordVector Averaging+TfIdf	90.67
WordVector Averaging+Document Vector	93.18
WordVector Averaging+Wiki+Document Vector	93.18
WordVector Averaging+Document Vector+RNNLM	93.70
WordVector Averaging+Wiki+Document Vector+RNNLM	93.57
WordVector Averaging+TfIdf+Document Vector	93.91
WordVector Averaging+Wiki+Document Vector+TfIdf	93.55
WordVector Averaging+TfIdf+Document Vector+RNNLM	94.19

Table 6.4: Comparison of results on IMDB Movie Review Dataset with Various Features




Figure 6.1: Accuracies of Different Classifiers with Average Word Vectors(IMDB).

Figure 6.1 gives an overview of accuracies of various classifiers using word vector averaging method on IMDB dataset. This clearly gives a green signal to us for using Linear SVM as our classier in all experiments.

Table 6.5 presents result of another experiment conducted on famous Amazon electronics review dataset(refer 3.2.3). Our vector averaging method alone has beaten previous best by 3.3% which was based on maximum entropy method.

Features	Accuracy
Dredze et al.(2008)	85.90
Max Entropy	83.79
WordVector Averaging (Our Method)	89.41
Composite Document Vector (Our Method)	92.17
Composite Document Vector + RNNLM	92.91

Table 6.5: Results on Amazon Electronics Review Dataset

6.2 Results on Hindi Datasets

Features	Accuracy(1)	Accuracy(2)
WordVector Averaging	78.0	79.62
WordVector+tf-idf	90.73	89.52
WordVector+tf-idf without stop words	91.14	89.97
Weighted WordVector	89.71	85.90
Weighted WordVector+tf-idf	92.89	90.30

Table 6.6: Accuracies for Product Review and Movie Review Datasets.

Table 6.6 represents the results using five different techniques for feature set construction. We see that there is a slight improvement in accuracy on both datasets once we remove stop-words but the major breakthrough occurs once we used weighted averaging technique for construction of document vectors from word vectors.

Experiment	Features	Accuracy
Subjective Lexicon (Bakliwal et al., 2012)	Simple Scoring	79.03
Hindi-SWN Baseline (Arora et al., 2013)	Adjective and Adverb presence	69.30
Word Vector with SVM (Our method)	tf-idf with word vector	91.14
Weighted Word Vector with SVM (Our method)	tf-idf+weighted word vector	92.89

Table 6.7: Comparison of Approaches: Product Review Dataset

Table 6.7 and 6.8 compares our best method with various other methods which have performed well using techniques such as *tf-idf*, subjective lexicon, etc.

Experiment	Features	Accuracy
In language using SVM (Joshi et al., 2010))	tf-idf	78.14
MT Based using SVM (Joshi et al., 2010)	tf-idf	65.96
Improved Hindi-SWN (Bakliwal et al., 2012)	Adjective and Adverb presence	79.0
WordVector Averaging	word vector	78.0
Word Vector with SVM (Our method)	tf-idf with word vector	89.97
Weighted Word Vector with SVM (Our method)	tf-idf+weighted word vector	90.30

Table 6.8: Comparison of Approaches: Movie Review Dataset

Table 6.9 shows the top few similar words for certain words from the corpus with cosine similarity as a distance metric. The words which have higher cosine similarity tend to be semantically and syntactically related.

अच्छा	खराब	भयानक
बहुत	निरासाजनक	भयन्कर
सुपर	कम्जोर	भीषण
केवल	नाजुक	भयावह
इतना	बदतर	अवसाद

Table 6.9: Some sentiment words and their neighbors

6.3 Odd One Out

We trained Skip-Gram model on latest English and Hindi Wikipedia dump and tried to analyze the quality of embeddings by finding the odd word amongst a list of words. For this task, we took found cosine similarity between each pair and found the word with lowest cosine similarity with every other. The results are highlighted below-

breakfast	cereal	lunch	dinner
eight	seven	owe	nine
shopping	math	reading	science

Table 6.10: Odd One Out in English

भारत	मुम्बई	रूस	चीन
लड़की	बेहन	मर्द	महिला
उद्योग	नेता	मंत्री	सरकार

Table 6.11: Odd One Out in Hindi

6.4 Similar Words

The trained skip-gram model was used for this task as well. We found the top few words which were closer in terms of cosine distance to the give word. The results are highlighted below-

Father	France	XBOX	scratched	megabits
grandfather	Germany	XBLA	scraped	gigabits
uncle	French	Xbox360	rubbed	kilobits
mother	Greece	SmartGlass	bruised	megabit
father-in-law	Netherlands	360/PS3	cracked	terabits
brother	Scotland	XBLA	discarded	MB/s
-	-	Qubed	shoved	Tbit/s
-	-	Kinect	tripped	-

Table 6.12: Top Few Similar words in English

भारत	व्यापार	ओबामा
प्रदेश	व्यवसाय	क्लिंटन
तिब्बत	पुनर्बीमा	बराक
देश	वाणिज्य	सीनेटर
आंध्रप्रदेश	बैंकिंग	राष्ट्रपति
लद्दाख	उद्योग	उम्मीदवार

Table 6.13: Top Few Similar words in Hindi

Chapter 7

Conclusion and Future Work

7.1 Conclusion

In this work we present an early experiment on the possibilities of distributional semantic models (word vectors) for low-resource, highly inflected languages such as Hindi. What is interesting is that our word vector averaging method along with tf-idf results in improvements of accuracy compared to existing state-of-the-art methods for sentiment analysis in Hindi (from 80.2% to 90.3% on IITB Movie Review Dataset).

We observe that pruning high-frequency stop words improves the accuracy by around 0.45%. This is most likely because such words tend to occur in most of the documents and don't contribute to sentiment. Similarly, words with very low frequency are noisy and can be pruned. For example, the word **फिल्म** occurs in 139/252 documents in Movie Review Dataset(55.16%) and has little effect on sentiment. Similarly words such as **सिद्धार्थ** occur in 2/252 documents in Movie Review Dataset(0.79%). These words don't provide much information.

We also observe that giving weights to word vectors while adding has resulted in a significant increase of accuracy. This seems much in accordance with expectation as pure addition ignores the importance of word which is directly related to its ratio of frequency in whole corpus and its frequency in a particular document. We have successfully tapped this area which most of the works have ignored.

We also see that when number of features accumulate to a large number than there are few redundant features present also. This brings a noise in the representation of the text in high dimension. We tried to reduce this noise by using feature variance technique called ANOVA-F and PCA and have succeeded as well. There is a large increase in accuracy(around 11%).

Before concluding, we return to the unexpectedly high improvement in accuracy achieved. One possibility we considered is that when the skip-grams are learned from the entire review corpus, it incorporates some knowledge of the test data. But this seems unlikely since the difference in including this vs not including it, is not too significant. The best explanation may be that the earlier methods, which were all in some sense based on a sentiWordnet, and at that one that was initially translated from English, were essentially very weak. This is also clear in an analysis from [1], which shows intern-annotator agreement on sentiment words are very poor (70%) - i.e. about 30% of these words have poor human agreement. Compared to this, the word vector model provides considerable power, especially as amplified by the tf-idf process. Thus, this also seems to underline the claim that distributional semantics is a topic worth exploring for Indian languages.

Our experiments on new dataset and existing datasets show that our method is competitive with existing methods including state-of-the-art. This concept of Enhanced Document Vectors can overcome the weaknesses of existing models which were either deficient in capturing syntactic or semantic properties of text. The ensemble of RNNLM and Enhanced Document Vector has beaten state-of-the-art by a significant margin and has opened this area for future research. These models have the advantage that they don't require parsing at any step neither do they require a lot of heavy pre-processing. These tasks require a lot of extra effort and they slow the progress a lot. We have made the code available at <http://github.com/pranjals16/thesis> for anyone to reproduce the result as well as extend this work further.

The challenge of obtaining a better corpus still remains with Hindi even though we

have released a much larger corpus of movie reviews for sentiment analysis and for other NLP tasks as well.

7.2 Future Work

Distributional semantics approaches remain relatively under-explored for Indian languages, and our results suggest that there may be substantial benefits to exploring these approaches for Indian languages. While this work has focused on sentiment classification, it may also improve a range of tasks from verbal analogy tests to ontology learning, as has been reported for other languages. For future work, we can explore various compositional models - a) weighted average - where weights are determined based on cosine distances in vector space; b) weighted multiplicative models. Another aspect we are considering is to incorporate multiple word vectors for the same surface token in cases of polysemy - this would directly be useful for word sense disambiguation. Identifying morphological variants would be another direction to explore for better accuracy. With regard to sentiment analysis, the idea of aspect-based models (or part-based sentiment analysis), which looks into constituents in a document and classify their sentiment polarity separately, remains to be explored in Hindi. Another point to note is that we are re-computing the word vectors for the two review corpora, which are extremely small. We may expect better performance with a larger sentiment corpus as is the case with English.

We see a possible scope of improvement if we could learn vectors for highly frequent n-grams as that would be a better representation than doing a vector composition of individual units. In English, our enhanced document vectors has given an indication that current representations are not sufficient to model sentences/documents and that ensembles could also prove useful in tasks such as sentiment analysis. Another interesting and open area is to look at *Region of Importance* in NLP where we filter out sentiment oriented sentences and phrases from a unfocused corpus which contains text from various domains.

7.3 Summary

- **Corpus:** This plays a very significant role in the process of sentiment prediction and has not been very widely tapped. Some corpus tend to be more focused towards sentiment while some contain a lot of information which is redundant as far as their sentiment score is concerned.
- **Features:** This parameter has been part of what we call in machine learning as *Feature Engineering* and has been very popular whenever a new algorithm appears in the domain of Machine Learning or Natural Language Processing. We have effectively tapped this area in our work to achieve excellent results.
- **Classifier:** It is the deciding factor on which whole classification task depends and selection of a suitable classifier according to the task is a crucial factor to look at. We have used Linear SVM in almost all experiments.
- **Ensemble:** It has improved our result significantly and is an interesting area to experiment with which has lacked much attention by many researchers in NLP. We have included a simple and effective ensemble method involving probability of classification in our work.

Bibliography

- [1] Akshat Bakliwal, Piyush Arora, and Vasudeva Varma. Hindi subjective lexicon: A lexical resource for hindi adjective polarity classification. In *Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC-2012)*. European Language Resources Association (ELRA), 2012.
- [2] Marco Baroni, Raffaella Bernardi, and Roberto Zamparelli. Frege in space: A program of compositional distributional semantics. *Linguistic Issues in Language Technology*, 9, 2014.
- [3] William Blacoe and Mirella Lapata. A comparison of vector-based representations for semantic composition. In *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pages 546–556. Association for Computational Linguistics, 2012.
- [4] Nivet Chirawichitchai. Emotion classification of thai text based using term weighting and machine learning techniques. In *Computer Science and Software Engineering (JCSSE), 2014 11th International Joint Conference on*, pages 91–96. IEEE, 2014.
- [5] Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. Natural language processing (almost) from scratch. *The Journal of Machine Learning Research*, 12:2493–2537, 2011.
- [6] Mark Dredze, Koby Crammer, and Fernando Pereira. Confidence-weighted linear classification. In *Proceedings of the 25th international conference on Machine learning*, pages 264–271. ACM, 2008.
- [7] Vasileios Hatzivassiloglou and Kathleen R McKeown. Predicting the semantic orientation of adjectives. In *Proceedings of the 35th annual meeting of the association for computational linguistics and eighth conference of the european chapter of the association for computational linguistics*, pages 174–181. Association for Computational Linguistics, 1997.
- [8] Eric H. Huang, Richard Socher, Christopher D. Manning, and Andrew Y. Ng. Improving Word Representations via Global Context and Multiple Word Prototypes. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2012.
- [9] Rie Johnson and Tong Zhang. Effective use of word order for text categorization with convolutional neural networks. *arXiv preprint arXiv:1412.1058*, 2014.

- [10] Thomas K Landauer and Susan T Dumais. A solution to plato's problem: The latent semantic analysis theory of acquisition, induction, and representation of knowledge. *Psychological review*, 104(2):211, 1997.
- [11] Thomas K Landauer, Darrell Laham, and Peter W Foltz. Automated scoring and annotation of essays with the intelligent essay assessor. *Automated essay scoring: A cross-disciplinary perspective*, pages 87–112, 2003.
- [12] Quoc V Le and Tomas Mikolov. Distributed representations of sentences and documents. *arXiv preprint arXiv:1405.4053*, 2014.
- [13] Yang Liu, Zhiyuan Liu, Tat-Seng Chua, and Maosong Sun. Topical word embeddings. 2015.
- [14] B. K. Matilal. *The Word and The World India's contribution to the study of language*. Oxford Paperback, Delhi, 2001.
- [15] Julian McAuley and Jure Leskovec. Hidden factors and hidden topics: understanding rating dimensions with review text. In *Proceedings of the 7th ACM conference on Recommender systems*, pages 165–172. ACM, 2013.
- [16] Nishantha Medagoda, Subana Shanmuganathan, and Jacqueline Whalley. A comparative analysis of opinion mining and sentiment classification in non-english languages. In *Advances in ICT for Emerging Regions (ICTer), 2013 International Conference on*, pages 144–148. IEEE, 2013.
- [17] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013.
- [18] Tomas Mikolov, Martin Karafiát, Lukas Burget, Jan Cernocký, and Sanjeev Khudanpur. Recurrent neural network based language model. In *INTER-SPEECH 2010, 11th Annual Conference of the International Speech Communication Association, Makuhari, Chiba, Japan, September 26-30, 2010*, pages 1045–1048, 2010.
- [19] Tomas Mikolov, Stefan Kombrink, Lukas Burget, Jan H Cernocky, and Sanjeev Khudanpur. Extensions of recurrent neural network language model. In *Acoustics, Speech and Signal Processing (ICASSP), 2011 IEEE International Conference on*, pages 5528–5531. IEEE, 2011.
- [20] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, pages 3111–3119, 2013.
- [21] Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. Linguistic regularities in continuous space word representations. In *HLT-NAACL*, pages 746–751, 2013.
- [22] Jeff Mitchell and Mirella Lapata. Vector-based models of semantic composition. In *In Proceedings of ACL-08: HLT*, pages 236–244, 2008.
- [23] Namita Mittal, Basant Agarwal, Garvit Chouhan, Nitin Bania, and Prateek Pareek. Sentiment analysis of hindi review based on negation and discourse relation. In *proceedings of International Joint Conference on Natural Language Processing*, pages 45–50, 2013.

- [24] Frederic Morin and Yoshua Bengio. Hierarchical probabilistic neural network language model. In *Proceedings of the international workshop on artificial intelligence and statistics*, pages 246–252. Citeseer, 2005.
- [25] Arvind Neelakantan, Jeevan Shankar, Alexandre Passos, and Andrew McCallum. Efficient nonparametric estimation of multiple embeddings per word in vector space. In *Proceedings of EMNLP*, 2014.
- [26] Bo Pang and Lillian Lee. Opinion mining and sentiment analysis. *Found. Trends Inf. Retr.*, 2(1-2):1–135, January 2008.
- [27] Jeffrey Pennington, Richard Socher, and Christopher Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543. Association for Computational Linguistics, 2014.
- [28] Gerard Salton and Christopher Buckley. Term-weighting approaches in automatic text retrieval. *Information processing & management*, 24(5):513–523, 1988.
- [29] Richa Sharma, Shweta Nigam, and Rekha Jain. Opinion mining in hindi language: A survey. *arXiv preprint arXiv:1404.4935*, 2014.
- [30] Pranjal Singh and Amitabha Mukerjee. Word vector models for hindi: A sentiment analysis evaluation. In *Regional Symposium on Natural Language Processing*. NLP Association of India, 2015.
- [31] Richard Socher, Alex Perelygin, Jean Y Wu, Jason Chuang, Christopher D Manning, Andrew Y Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the conference on empirical methods in natural language processing (EMNLP)*, volume 1631, page 1642. Citeseer, 2013.
- [32] Joseph Turian, Lev Ratinov, and Yoshua Bengio. Word representations: A simple and general method for semi-supervised learning. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics, ACL ’10*, pages 384–394, Stroudsburg, PA, USA, 2010. Association for Computational Linguistics.
- [33] Peter D Turney, Patrick Pantel, et al. From frequency to meaning: Vector space models of semantics. *Journal of artificial intelligence research*, 37(1):141–188, 2010.
- [34] L.J.P. van der Maaten and G.E. Hinton. Visualizing high-dimensional data using t-sne. 2008.
- [35] Gang Wang, Jianshan Sun, Jian Ma, Kaiquan Xu, and Jibao Gu. Sentiment classification: The contribution of ensemble learning. *Decision support systems*, 57:77–93, 2014.
- [36] Wikipedia. Principal component analysis, 2015. [Online; accessed 10-April-2015].

- [37] Will Y Zou, Richard Socher, Daniel M Cer, and Christopher D Manning. Bilingual word embeddings for phrase-based machine translation. In *EMNLP*, pages 1393–1398, 2013.